

profiling-plots

December 7, 2023

1 Introduction

MESH has been compiled in serial mode with `netcdf` capability against the following libraries on [Digital Research Alliance of Canada's \(DRA\) Graham](#): 1. `intel/2020.1.217` Fortran `ifort` compiler, 2. `netcdf-fortran/4.5.2`, 3. `netcdf/4.7.4`, and 4. `hdf5/1.10.0`.

Two versions of MESH has been compiled, namely `r1773` (commit [75d48bd](#)) and `r1860ME` (commit [52c7367](#)). Both have been profiled using Intel's `vtune/2020.1` profiler on DRA Graham.

2 Issues encountered

1. `r1860ME` (latest commit [52c7367](#)) compiles successfully but raises the following runtime error running any MESH model setup if the `ns` value is assigned to `BASINAVGWBFILFLAG` flag of the `MESH_input_run_options.ini` file,

```
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4215 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
WARNING: NON-CONVERGENCE AT POINT AT X,Y: 4873 1
forrtl: severe (174): SIGSEGV, segmentation fault occurred
Image                PC                Routine                Line                Source
sa_mesh_1860me_se     0000000000B51A3A    Unknown                Unknown              Unknown
libpthread-2.30.s     00002AEA3960A0F0    Unknown                Unknown              Unknown
sa_mesh_1860me_se     0000000000AA9BF6    save_basin_output      1376                save_basin_output.f90
sa_mesh_1860me_se     0000000000AA1DB8    save_basin_output      766                  save_basin_output.f90
sa_mesh_1860me_se     0000000000B2191E    MAIN__                  1031                MESH_driver.f90
sa_mesh_1860me_se     000000000040CC12    Unknown                Unknown              Unknown
libc-2.30.so          00002AEA3963CE1B    __libc_start_main       Unknown              Unknown
sa_mesh_1860me_se     000000000040CB2A    Unknown                Unknown              Unknown
```

2. [r1860ME](#) (older commit [b7f23d1](#)) does not exhibit any issues with the `ns` flag in MESH model setups and runs successfully.
3. [r1860ME](#) (latest commit [52c7367](#)) compiles successfully with `gfortran/9.3.0` compiler (with relevant dependencies), but echoes the following runtime error running any MESH model setup:

```
RUNCLASS36 is active.
ICEBAL_FREEZE_THRESHOLD (FREZTH) override is ACTIVE.
Uniform value:      -2.000000
ICEBAL_SWE_LIMIT (SWELIM) override is ACTIVE.
Uniform value:       100.0000
ICEBAL_SNOW_DENSITY_LIMIT (SNDENLIM) override is ACTIVE.
Uniform value:       900.0000
BASEFLOW component is ACTIVE.
  BASEFLOWFLAG  wf_lzs grid hf=60
                pwr_iak   2.203000
                flz_iak   0.2800000E-04
```

Program received signal SIGSEGV: Segmentation fault - invalid memory reference.

```
Backtrace for this error:
#0  0x2b8dd7c99730 in ???
#1  0x2b8dd7c988d5 in ???
#2  0x2b8dd813497f in ???
#3  0x605e24 in __output_variables_MOD_output_variables_group_update_ts
at ./Driver/MESH_Driver/output_variables.f90:1389
#4  0x602d21 in __output_variables_MOD_output_variables_update_ts
at ./Driver/MESH_Driver/output_variables.f90:2078
#5  0x5ff384 in __output_variables_MOD_output_variables_update
at ./Driver/MESH_Driver/output_variables.f90:2530
#6  0x91e471 in runmesh
at ./Driver/MESH_Driver/MESH_driver.f90:847
#7  0x922ab2 in main
at ./Driver/MESH_Driver/MESH_driver.f90:97
Segmentation fault
```

3 Profiling

```
[1]: # import libraries
import pandas as pd
import matplotlib.pyplot as plt

import os
import time
```

3.1 Total CPU time

Now, let's check the bottlenecks of MESH r1860ME version and its differences with r1773. Only 1 year of the Fraser River Basin setup was run in serial mode by each version.

```
[2]: # read the profiling results
r1773 = pd.read_csv('./profiling-results/1773.csv')
r1860me = pd.read_csv('./profiling-results/1860me.csv')
```

Determining total CPU time used by each version (not considering the wall time):

```
r1773 used 04:58:08
r1860me used 05:37:05
```

3.2 Function call differences between r1860ME and r1773

Determining functions that are called/available in r1860ME but not in r1773 plus see how much time they have used the CPU:

```
[4]: diffs = r1860me[~r1860me.iloc[:,0].isin(r1773.iloc[:,0])].sort_values(by='CPU_
↳Time', axis='rows', ascending=False)
```

The top 6 new function calls of r1860ME with considerable computation time are listed in the following:

```
[5]: Source Function / Function / Call Stack    CPU Time    Module \
3      copy_field_scalar_to_scalar    1040.2800    [Unknown]
14      map_field_2d    576.4670    [Unknown]
19      func@0x151900    531.8390    [Unknown]
26      map_field_to_ranked_output    193.3040    [Unknown]
55      func@0x40c720    28.0808    [Unknown]
56      read_frame_from_file    26.1351    [Unknown]
```

| | Function (Full) | Source File | Start Address |
|----|-----------------------------|---------------------|---------------|
| 3 | copy_field_scalar_to_scalar | field_utilities.f90 | 0 |
| 14 | map_field_2d | field_utilities.f90 | 0 |
| 19 | func@0x151900 | [Unknown] | 0 |
| 26 | map_field_to_ranked_output | mesh_io.F90 | 0 |
| 55 | func@0x40c720 | [Unknown] | 0 |
| 56 | read_frame_from_file | mesh_io.F90 | 0 |

From the profiling analysis, the top 10 function calls introduced newly in r1860ME can explain the time difference of ~40 minutes:

New function calls in r1860ME are responsible for 00:40:12

It seems that the following stacks/files added in r1860ME are mostly in charge of the differences:

Function calls:
field_utilities.f90,
mesh_io.F90,

```
sa_mesh_run_within_grid.f90,
variable_maps.f90
```

3.3 Analysis of common function calls in r1860ME and r1773

```
[11]: r1860me = r1860me.set_index('Source Function / Function / Call Stack').loc[:, 'CPU Time'].groupby('Source Function / Function / Call Stack').sum().sort_values(ascending=False)

[12]: r1773 = r1773.set_index('Source Function / Function / Call Stack').loc[:, 'CPU Time'].groupby(['Source Function / Function / Call Stack']).sum().sort_values(ascending=False)

[13]: all = pd.merge(r1860me,
                    r1773,
                    left_index=True,
                    right_index=True,
                    how='inner',
                    suffixes=('_r1860me', '_r1773'),
                    )

[14]: total_common_delay = (all['CPU Time_r1860me'] - all['CPU Time_r1773']).sum()
```

For the common functions calls shared between the two versions, the r1860ME is performing quite close to the r1773 version.

Overall common function calls in r1860ME are about 00:00:59 slower than r1773.

And, here is a list of top 10 common functions and their processing times in the two version:

| [16]: | CPU Time_r1860me | CPU Time_r1773 |
|---|------------------|----------------|
| Source Function / Function / Call Stack | | |
| classw | 1396.490 | 1244.420 |
| grdran | 1083.770 | 990.855 |
| runclass36_within_tile | 1045.140 | 1183.150 |
| wflow | 955.822 | 956.044 |
| tmcalc | 886.655 | 890.895 |
| __libm_powf_l9 | 791.024 | 671.511 |
| aprep | 773.889 | 817.789 |
| flxsurfz | 676.447 | 692.902 |
| watrof | 663.026 | 664.725 |
| tsolvc | 658.972 | 704.973 |

3.4 Overall module contributions in each version

```
[18]: td_1773 = pd.read_csv('./profiling-results/top-down_1773.csv')
      td_1860me = pd.read_csv('./profiling-results/top-down_1860.csv')
```

Here, you can see the percentage of each module within MESH r1773 taking over the total computation:

[57]: Function Stack

| | |
|----------------------------|------------|
| runmesh | 100.000000 |
| run_within_tile | 94.030769 |
| run_between_grid | 3.043447 |
| run_within_grid | 2.192695 |
| climate_module_update_data | 0.648610 |
| output_variables_reset | 0.068582 |
| run_save_basin_output | 0.007378 |
| read_initial_inputs | 0.003298 |
| run_within_tile_init | 0.001453 |
| run_within_tile_finalize | 0.001342 |

Name: CPU Time: Total, dtype: float64

And, also for r1860ME:

[59]: Function Stack

| | |
|--------------------------|------------|
| runmesh | 100.000000 |
| run_within_tile | 82.421842 |
| read_input_forcing_frame | 12.110546 |
| run_within_grid | 2.706731 |
| run_between_grid | 2.698741 |
| output_variables_reset | 0.046276 |
| run_save_basin_output | 0.006081 |
| read_initial_inputs | 0.003263 |
| open_input_forcing_files | 0.002670 |
| run_within_tile_finalize | 0.001187 |

Name: CPU Time: Total, dtype: float64

4 Final notes

It seems that the I/O processes are causing newer version to take more time to complete, and therefore, there is a difference in the computation time.

There are many other details hidden in the profiling results, they could be shared with you. Let me know if you need to have access to them.

Profiling results are done using vtune/2020.1 on DRA Graham